

# Diffusion Models and SDEs

## Lecture 2: SDEs and how to manipulate them

Time Reversal and  
Score-based modelling

Probability Flow ODE and  
Flow Matching

Conditioning SDEs via  
Doob's h-transform

# Score-based generative modelling

*Training:* learn score  $\mathbf{s}_t(x) := \nabla \log p_t(x)$  by

1. Sample from target  $x_1 \sim p_1$
2. Run "forward" SDE to "noise"  $x_1$  until it becomes "simple"  $x_0 \sim p_0$
3. Minimize some loss  $\propto \|\hat{s}_t(x) - \mathbf{s}_t(x)\|^2$

*Inference:* sample using "reverse" SDE or prob-flow ODE

---

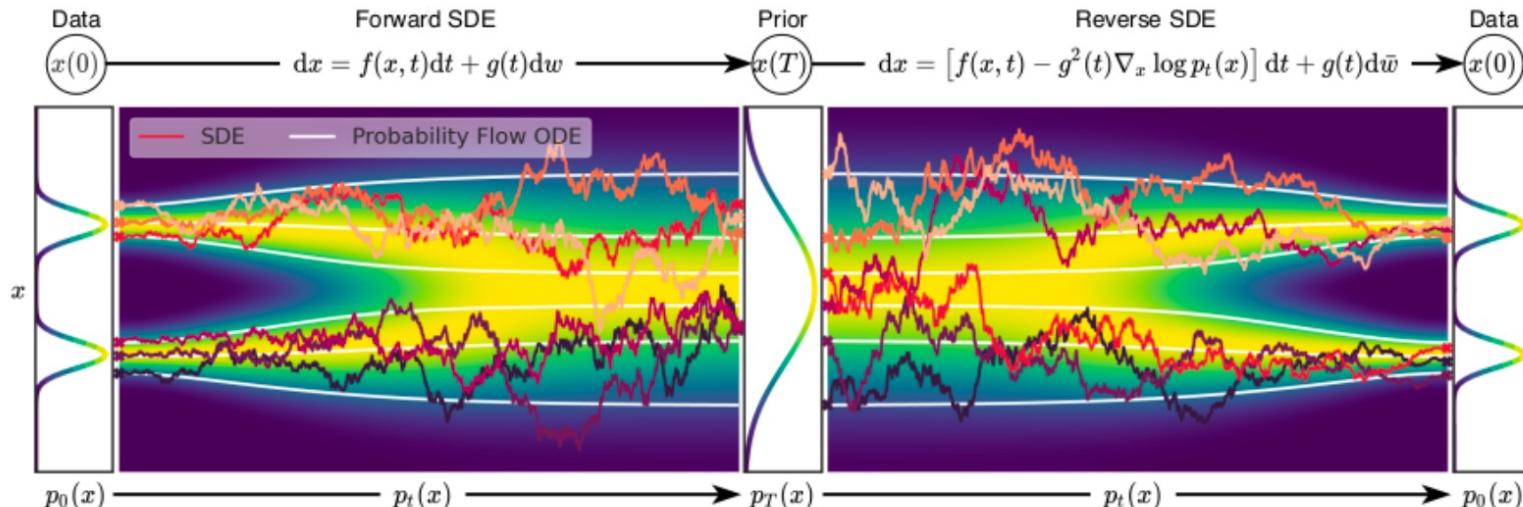


Figure 2 from Song et al. (2020)

# Time Reversal - Chain Rule

A discrete time “heuristic” sketch

Via the chain rule we can decompose the joint in either direction,

$$p_{t|\delta}(x|y)p_{\delta}(y) = p_{\delta|t}(y|x)p_t(x)$$

Now consider an EM approx transition density, for the forward kernel:

$$p_{\delta|t}(y|x) = \mathcal{N}(y|x + f^+(x)\delta, \delta\sigma^2)$$

$$p_{t|\delta}(x|y) = ?$$

# Time Reversal - Chain Rule

A discrete time “heuristic” sketch

Via the chain rule we can decompose the joint in either direction,

$$p_{t|\delta}(x|y)p_{\delta}(y) = p_{\delta|t}(y|x)p_t(x)$$

Now consider an EM approx transition density, for the forward kernel:

$$p_{\delta|t}(y|x) = \mathcal{N}(y|x + f^+(x)\delta, \delta\sigma^2)$$

$$p_{t|\delta}(x|y) = p_{\delta|t}(y|x) \frac{p_t(x)}{p_{\delta}(y)}$$

# Time Reversal - Chain Rule

A discrete time “heuristic” sketch

Via Taylors Theorem we can expand time t marginal around y:

$$p_{t|\delta}(x|y) = p_{t+\delta|t}(y|x) \frac{p_t(y) e^{(x-y)^\top \nabla_y \ln p_t(y) + \mathcal{O}(\delta^2)}}{p_{t+\delta}(y)}$$

Assuming  $|\ln p_t(x) - \ln p_s(x)| = \mathcal{O}(|t-s|^2)$

$$p_{t|\delta}(x|y) = p_{t+\delta|t}(y|x) e^{(x-y)^\top \nabla_y \ln p_t(y) + \mathcal{O}(\delta^2)}$$

# Time Reversal - Chain Rule

A discrete time “heuristic” sketch

Regrouping and completing the square:

$$p_{t|\delta}(x|y) = \frac{e^{-\frac{\|x - (y - f^+(y)\delta + \sigma^2 \nabla_y \ln p_t(y)\delta)\|^2}{\sigma^2 \delta}} + \mathcal{O}(\delta^2)}{\sqrt{2\pi} \delta^{d/2} \sigma^d}$$

Which corresponds to the Euler Maruyama discretization of the following SDE (seem familiar ?):

$$dX_t = (-f^+(X_t, T-t) + \sigma^2 \nabla_{X_t} \ln p_{T-t}(X_t)) dt + \sigma dW_t$$

# Time Reversal - Chain Rule

A discrete time “heuristic” sketch

Inspecting the relationship between the drifts yields Nelsons duality formula:

$$f^-(x, t) + f^+(x, T - t) = \sigma^2 \nabla_x \ln p_{T-t}(x)$$



# Time Reversal - Chain Rule

A discrete time “heuristic” sketch

Inspecting the relationship between the drifts yields Nelsons duality formula:

$$f^-(x, t) + f^+(x, T - t) = \sigma^2 \nabla_x \ln p_{T-t}(x)$$

Looks slightly different to Song et al. 2021, why ?

# Time Reversal - Chain Rule

## Nelsons Relation – Semantics Clarification

Looks slightly different to Song et al. 2020, why ?

Due to 2 equivalent ways of representing time reversals:

$$dY_t = f^+(Y_t, t)dt + \sigma dW_t$$

### Forward SDE (e.g. De Bortoli 2021)

- Travels forward in time

$$dX_t = f^-(X_t, t)dt + \sigma dW_t$$
$$f^-(x, t) + f^+(x, T-t) = \sigma^2 \nabla_x \ln p_{T-t}(x)$$

- Flips / No longer the same joint

$$\text{Law}(x_t)_{t=0}^T = \text{Law}(y_{T-t})_{t=0}^T$$

### Backwards SDE (e.g. Song 2021)

- Travels Backwards in time

$$dX_t^- = f^-(X_t^-, t)dt + \sigma dW_t^-$$
$$f^-(x, t) - f^+(x, t) = \sigma^2 \nabla_x \ln p_t(x)$$

- Encodes the same joint

$$\text{Law}(x_t)_{t=0}^T = \text{Law}(y_t)_{t=0}^T$$

# Time Reversal – Generative Modelling

**Time reversing VP-SDE / OU Process [Song 2021, De Bortoli 2021]**

Consider the time homogenous VP-SDE (OU Process):

$$X_0 \sim p_{\text{data}}$$
$$dX_t = -\beta X_t dt + \sqrt{2\beta} dW_t$$

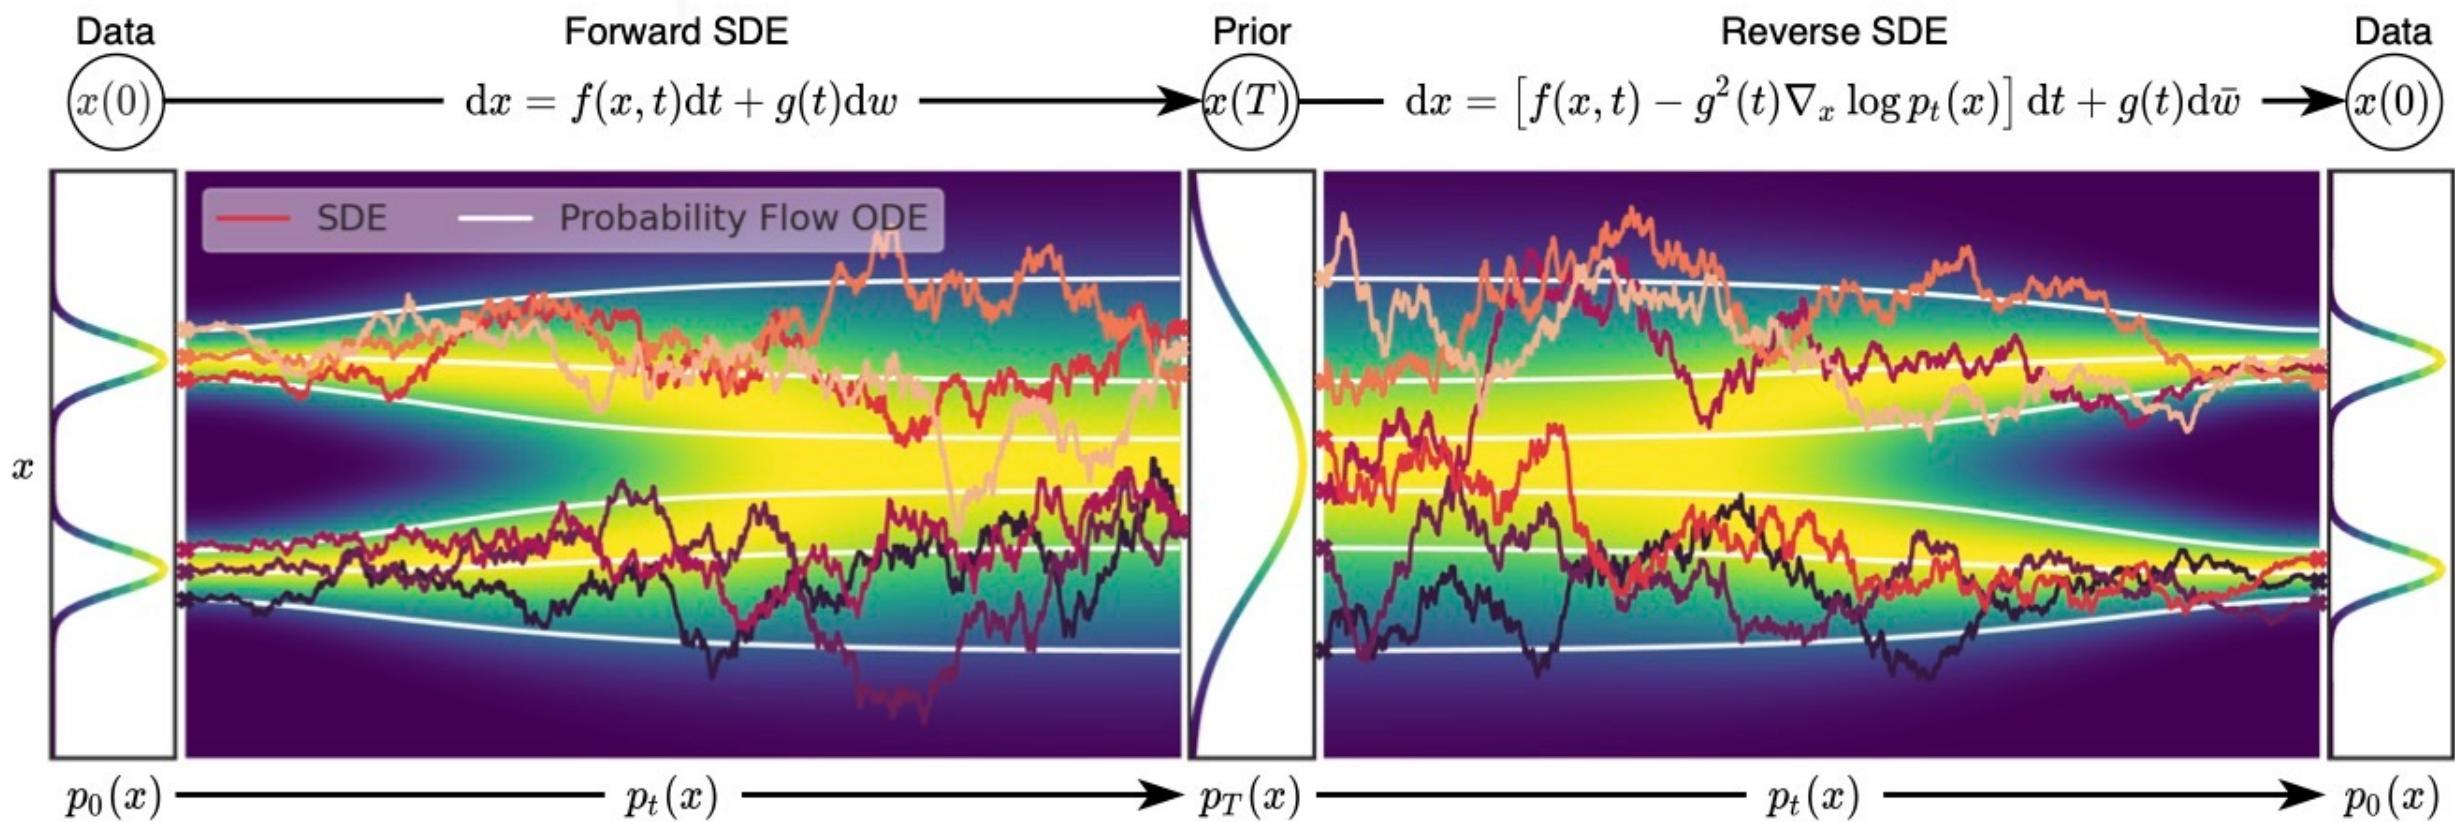
Then its time reversal  $(Y_t)_{t=0}^T \stackrel{d}{=} (X_{T-t})_{t=0}^T$  satisfies the score SDE [Song 2021]:

$$Y_0 \sim p_T \approx \mathcal{N}(0, I)$$

$$dY_t = (\alpha Y_t + 2\alpha \nabla_{Y_t} \ln p_{T-t}(Y_t)) dt + \sqrt{2\alpha} dB_t$$

Where  $Y_T \sim p_{\text{data}}$ , thus we could instead sample approximately  $Y_0 \sim \mathcal{N}(0, I)$  and have  $\text{Law } Y_T \approx p_{\text{data}}$  following the mixing rate of the OU [De Bortoli 2021]

# Probability flow ODE



# Probability Flow ODE

Definition: Every stochastic process described by an SDE has a corresponding deterministic process described by an ODE that has the same marginal probability densities  $\{p_t(x)\}_{t=0}^T$ . This process is called the *probability flow ODE*. For a general SDE of the form  $dX_t = \mu(X_t, t)dt + \sigma(X_t, t)dW_t$ , the corresponding ODE is given by

$$dX_t = \left[ \mu(X_t, t) - \frac{1}{2} \nabla_x [\sigma(X_t, t) \sigma(X_t, t)^T] - \frac{1}{2} \sigma(X_t, t) \sigma(X_t, t)^T \nabla_x \log p_t(x) \right] dt$$

# Probability Flow ODE

$$dX_t = \left[ \mu(X_t, t) - \frac{1}{2} \nabla_x [\sigma(X_t, t) \sigma(X_t, t)^T] - \frac{1}{2} \sigma(X_t, t) \sigma(X_t, t)^T \nabla_x \log p_t(x) \right] dt$$

$$dX_t = [\mu(X_t, t) - \frac{1}{2} \sigma^2(t) \nabla_x \log p_t(X_t)] dt$$

# Probability Flow ODE

FPE

$$\frac{\partial}{\partial t} p_t(x) = -\frac{\partial}{\partial x} [\mu(x, t)p_t(x)] + \frac{\partial}{\partial x} \frac{\partial}{\partial x} \left[ \frac{1}{2} \sigma(x, t) \sigma(x, t)^T p_t(x) \right]$$

Product Rule

$$\frac{\partial}{\partial t} p_t(x) = -\frac{\partial}{\partial x} [\mu(x, t)p_t(x)] + \frac{1}{2} \frac{\partial}{\partial x} \left[ \nabla_x [\sigma(x, t) \sigma(x, t)^T] p_t(x) \right] + \frac{1}{2} \frac{\partial}{\partial x} \left[ \sigma(x, t) \sigma(x, t)^T \frac{\partial p_t(x)}{\partial x} \right]$$

Log Der. Trick

$$\frac{\partial}{\partial t} p_t(x) = -\frac{\partial}{\partial x} [\mu(x, t)p_t(x)] + \frac{1}{2} \frac{\partial}{\partial x} \left[ \nabla_x [\sigma(x, t) \sigma(x, t)^T] p_t(x) \right] + \frac{1}{2} \frac{\partial}{\partial x} \left[ \sigma(x, t) \sigma(x, t)^T p_t(x) \nabla_x \log p_t(x) \right]$$

# Probability Flow ODE

$$\frac{\partial}{\partial t} p_t(x) = -\frac{\partial}{\partial x} [\mu(x, t)p_t(x)] + \frac{1}{2} \frac{\partial}{\partial x} \left[ \nabla_x [\sigma(x, t)\sigma(x, t)^T] p_t(x) \right] + \frac{1}{2} \frac{\partial}{\partial x} \left[ \sigma(x, t)\sigma(x, t)^T p_t(x) \nabla_x \log p_t(x) \right]$$

Pull der.  
And  $p(x)$   
out

$$\frac{\partial}{\partial t} p_t(x) = -\frac{\partial}{\partial x} \left[ \mu(x, t) - \frac{1}{2} \nabla_x [\sigma(x, t)\sigma(x, t)^T] - \frac{1}{2} \sigma(x, t)\sigma(x, t)^T \nabla_x [\log p_t(x)] \right]$$

$$\frac{\partial}{\partial t} p_t(x) = -\frac{\partial}{\partial x} \tilde{\mu}(x, t)p_t(x)$$

# Generative Modelling

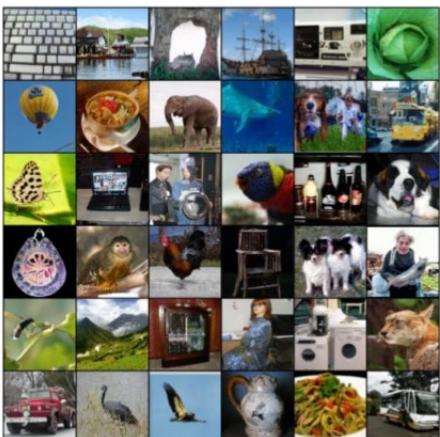
Given samples  $x_1, \dots, x_n$  with  $x \sim p$

Infer  $\hat{p}$  such that  $\hat{p} \approx p$ .

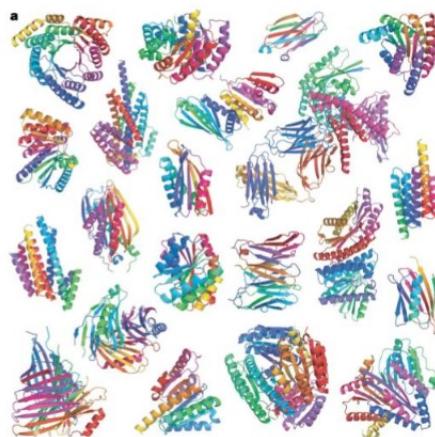
## Tasks

Evaluate the likelihood of new data  $\hat{p}(x)$

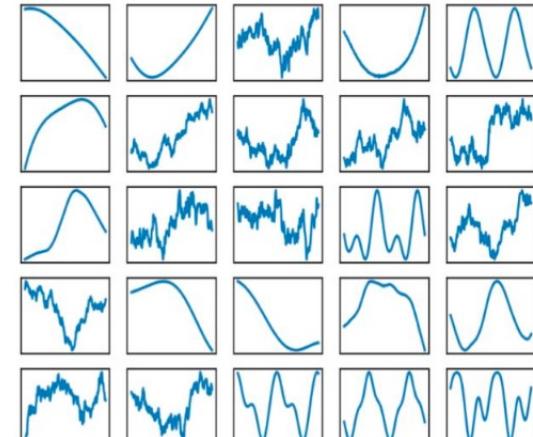
Sample  $x \sim \hat{p}$



Images



Proteins



Functions

# Generative Modelling

$$\max_{\theta} \mathbb{E}_{x \sim p_{\text{data}}} \log p_{\theta}(x)$$

# Generative Modelling

$$\max_{\theta} \mathbb{E}_{x \sim p_{\text{data}}} \log p_{\theta}(x)$$

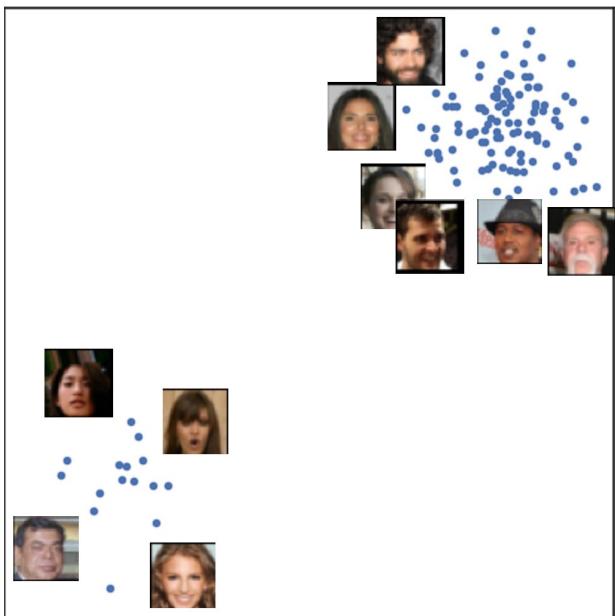
$$\begin{aligned}s_{\theta}(x) &= \nabla_x \log p_{\theta}(x) \\&= \nabla_x \log \left( \frac{f_{\theta}(x)}{Z_{\theta}} \right) \\&= \nabla_x \log f_{\theta}(x) - \underbrace{\nabla_x \log Z_{\theta}}_{=0} \\&= \nabla_x \log f_{\theta}(x)\end{aligned}$$

# How to train and simulate?

$$\mathbb{E}_{p(\mathbf{x})}[\|\nabla_{\mathbf{x}} \log p(\mathbf{x}) - \mathbf{s}_\theta(\mathbf{x})\|_2^2]$$

$$\mathbb{E}_{p(\mathbf{x})}[\|\nabla_{\mathbf{x}} \log p(\mathbf{x}) - \mathbf{s}_\theta(\mathbf{x})\|_2^2] = \int p(\mathbf{x}) \|\nabla_{\mathbf{x}} \log p(\mathbf{x}) - \mathbf{s}_\theta(\mathbf{x})\|_2^2 d\mathbf{x}.$$

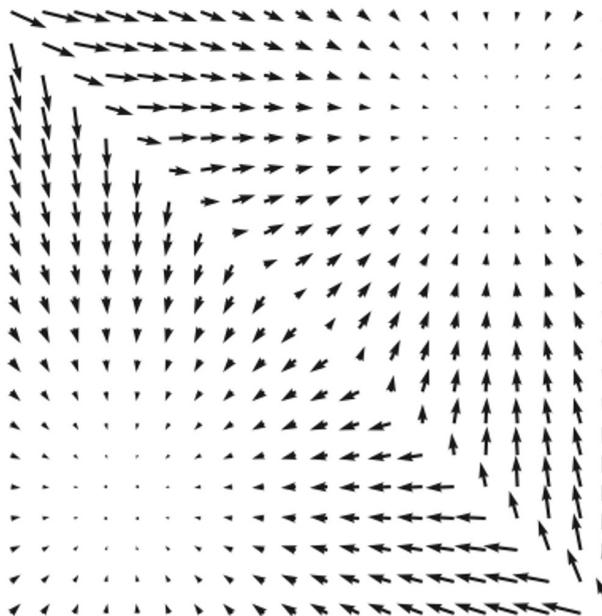
# Naïve Score Matching



Data samples

$$\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N\} \stackrel{\text{i.i.d.}}{\sim} p(\mathbf{x})$$

score  
matching



Scores

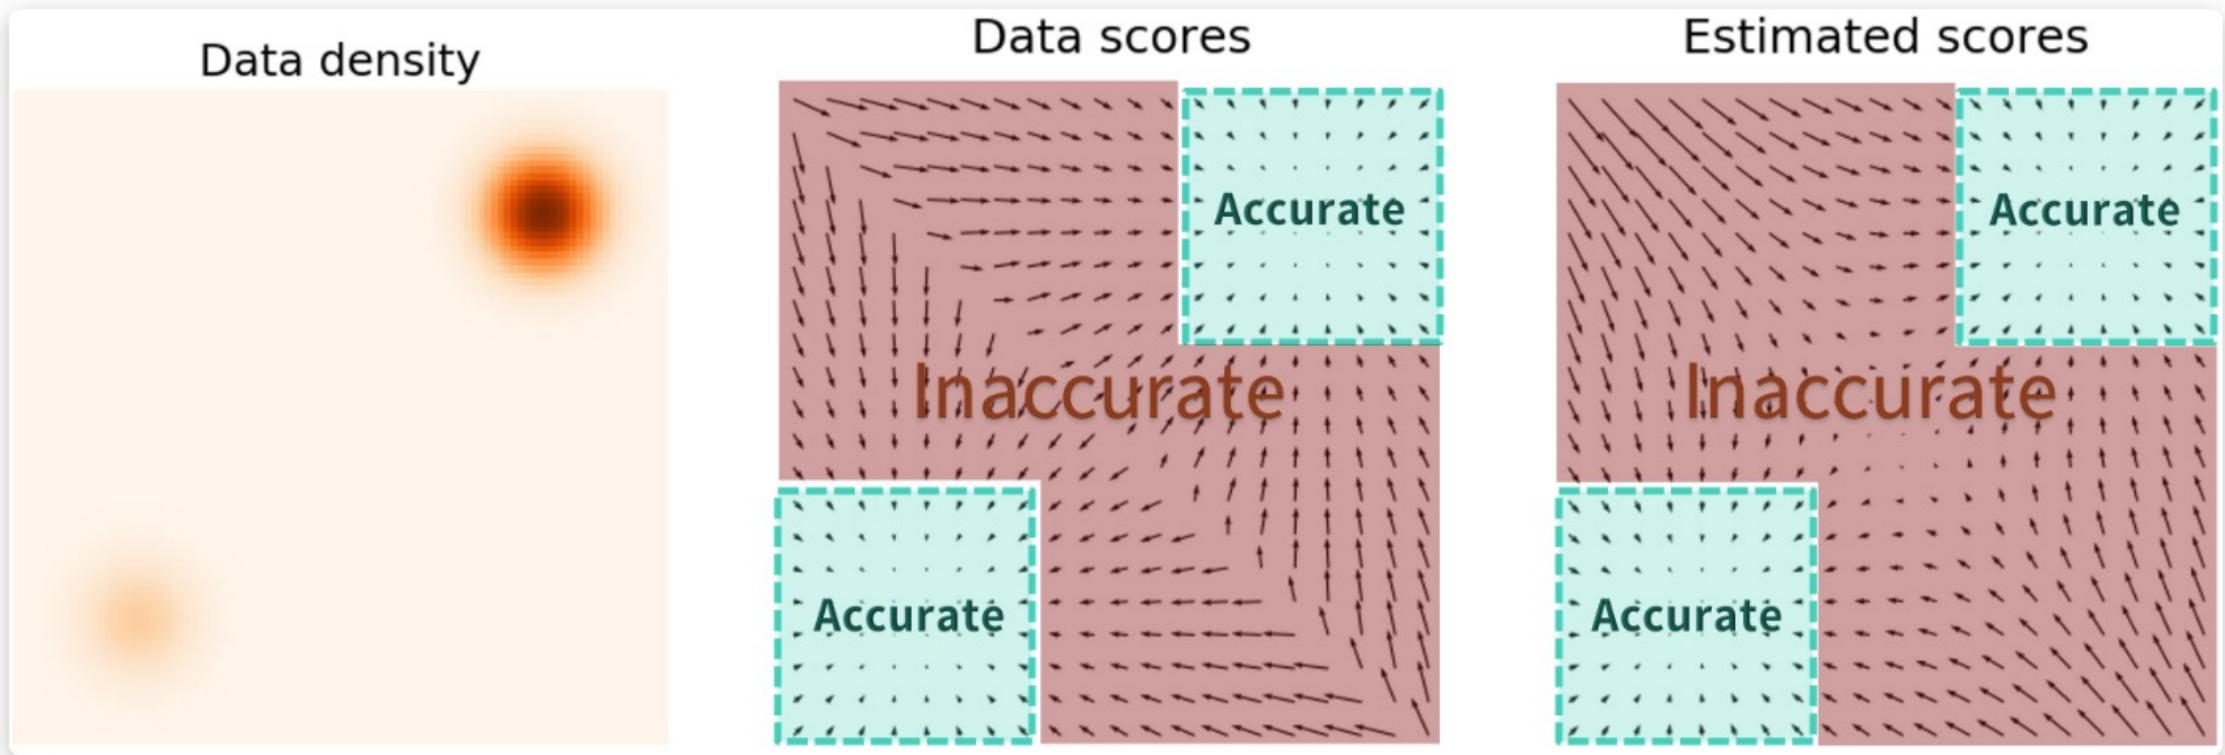
$$\mathbf{s}_\theta(\mathbf{x}) \approx \nabla_{\mathbf{x}} \log p(\mathbf{x})$$

Langevin  
dynamics



New samples

# Naïve Score Matching



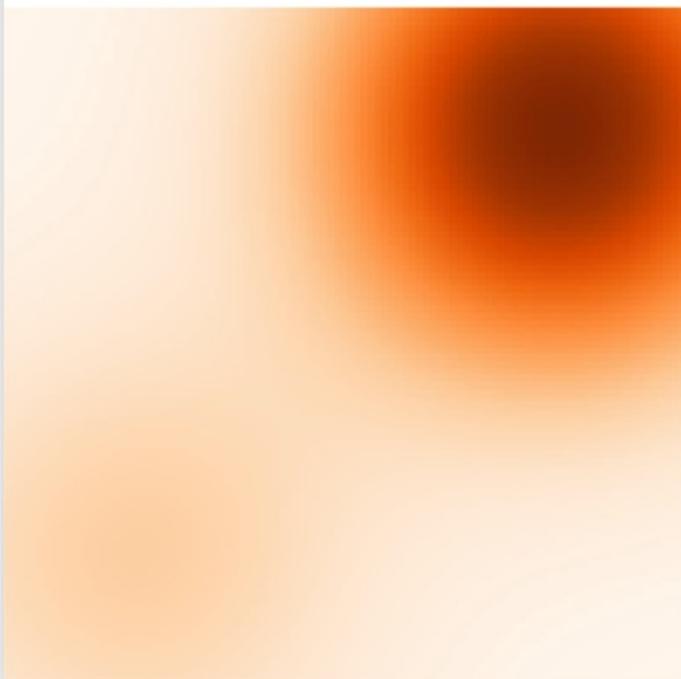
# How to train and simulate?

$$\mathbb{E}_{p(\mathbf{x})}[\|\nabla_{\mathbf{x}} \log p(\mathbf{x}) - \mathbf{s}_\theta(\mathbf{x})\|_2^2]$$

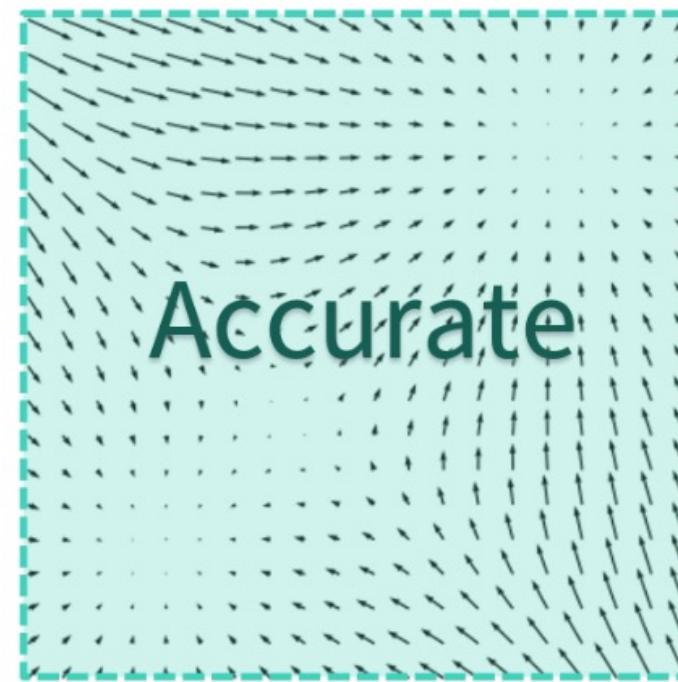
$$\mathbb{E}_{t \in \mathcal{U}(0,T)} \mathbb{E}_{p_t(\mathbf{x})} [\lambda(t) \|\nabla_{\mathbf{x}} \log p_t(\mathbf{x}) - \mathbf{s}_\theta(\mathbf{x}, t)\|_2^2],$$

# Denoising Score Matching

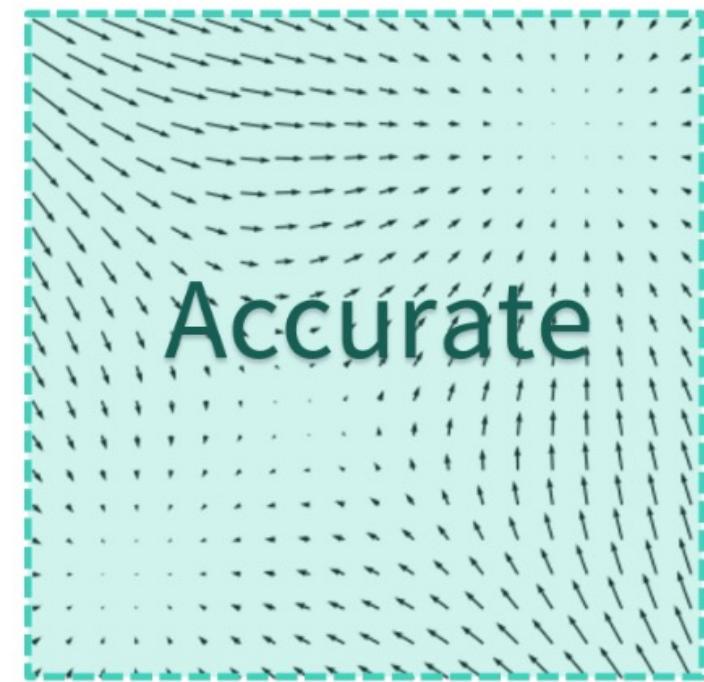
Perturbed density



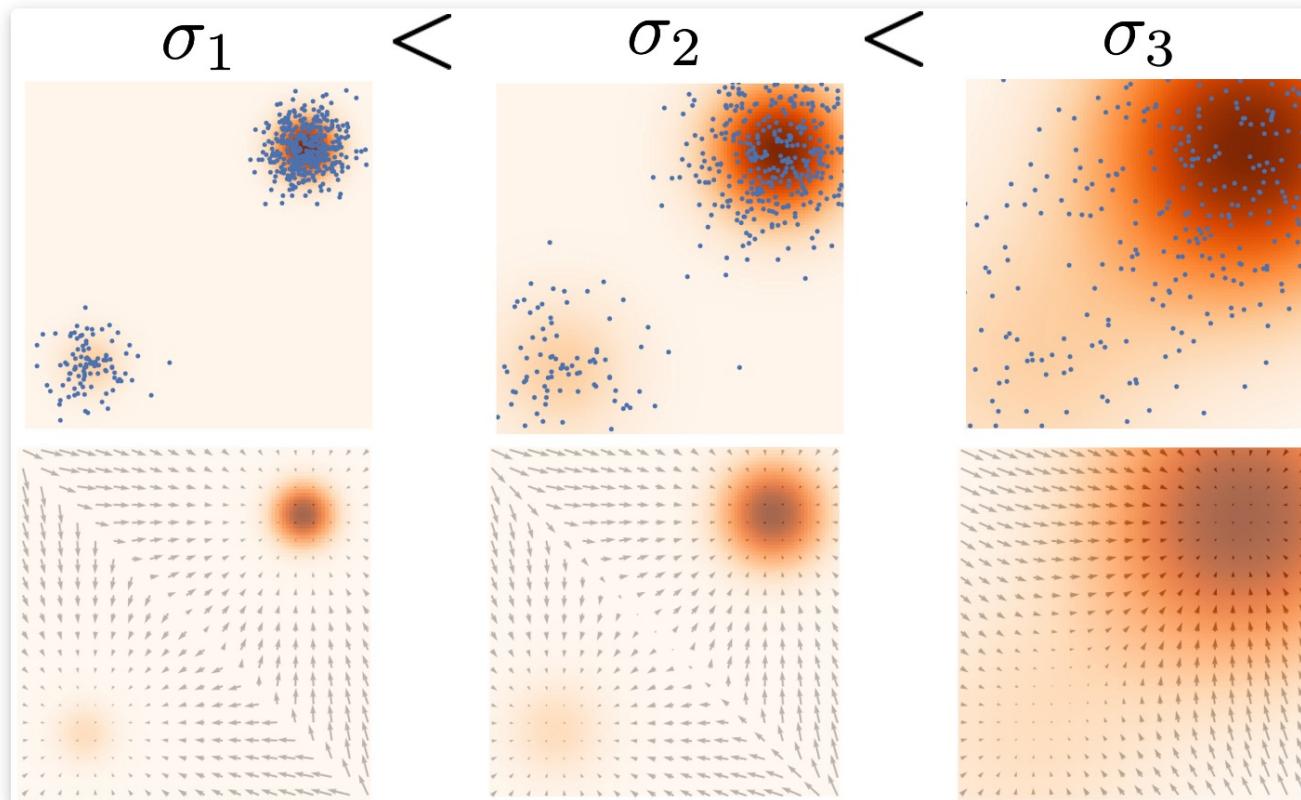
Perturbed scores



Estimated scores



# Denoising Score Matching



We apply multiple scales of Gaussian noise to perturb the data distribution (**first row**), and jointly estimate the score functions for all of them (**second row**).



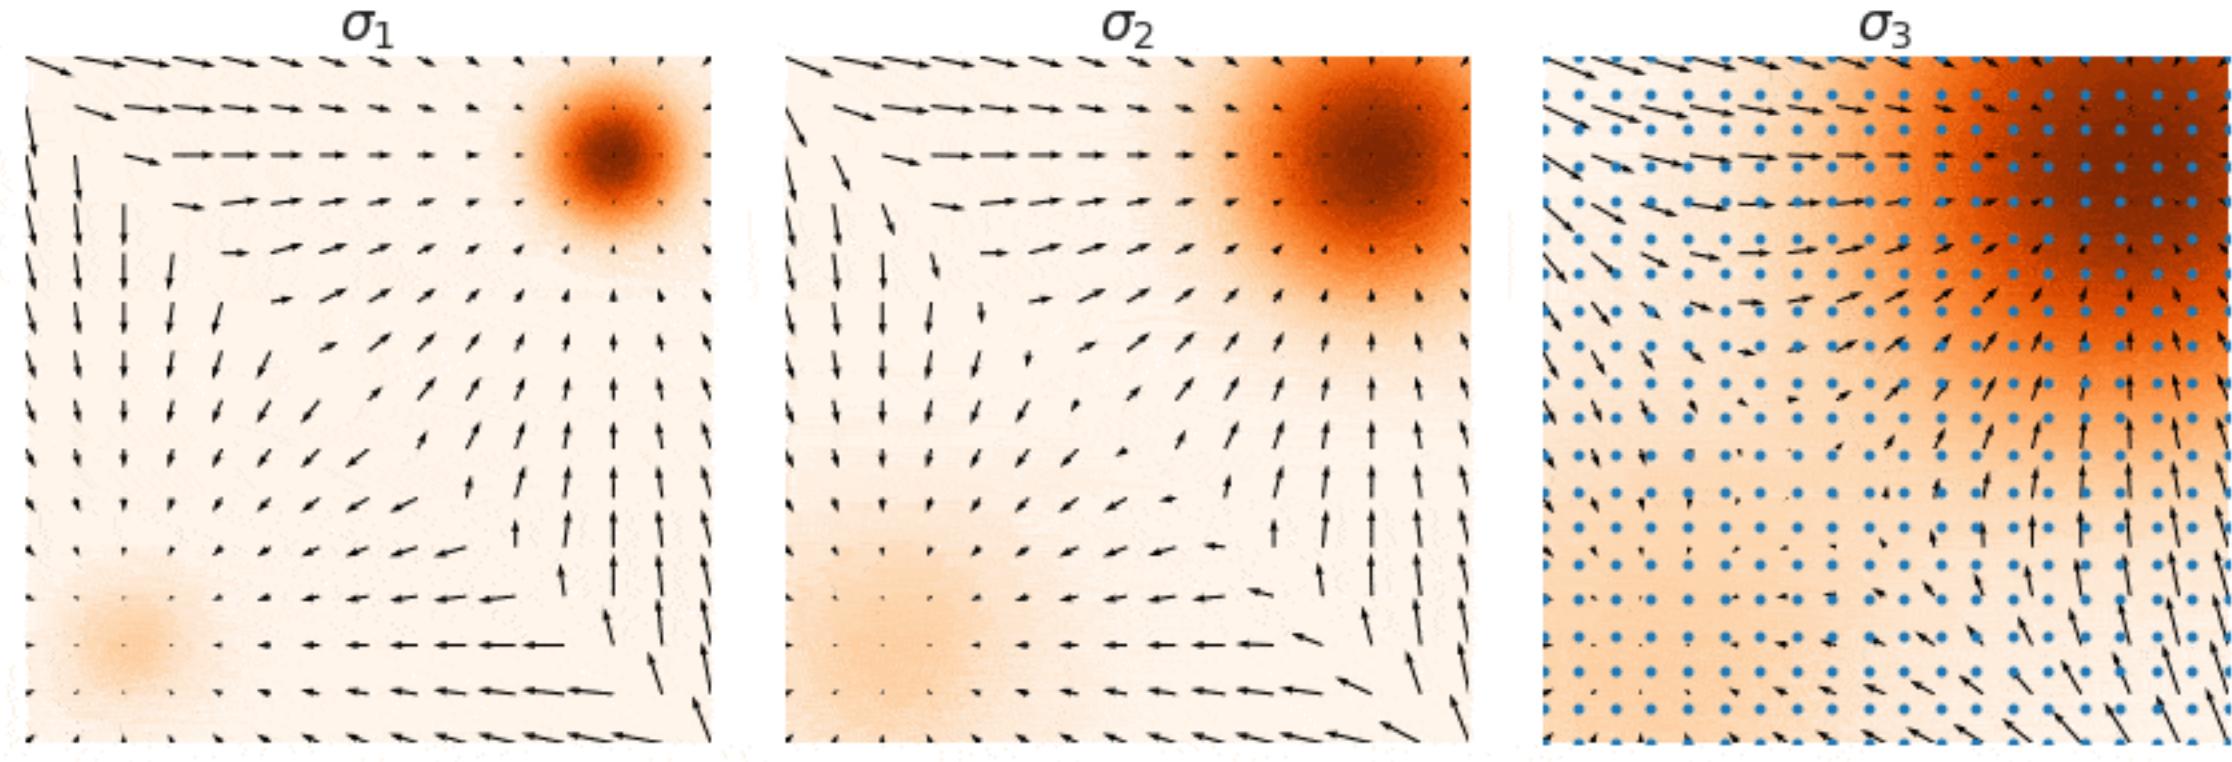
# How to train and simulate?

$$\mathbb{E}_{p(\mathbf{x})}[\|\nabla_{\mathbf{x}} \log p(\mathbf{x}) - \mathbf{s}_\theta(\mathbf{x})\|_2^2]$$

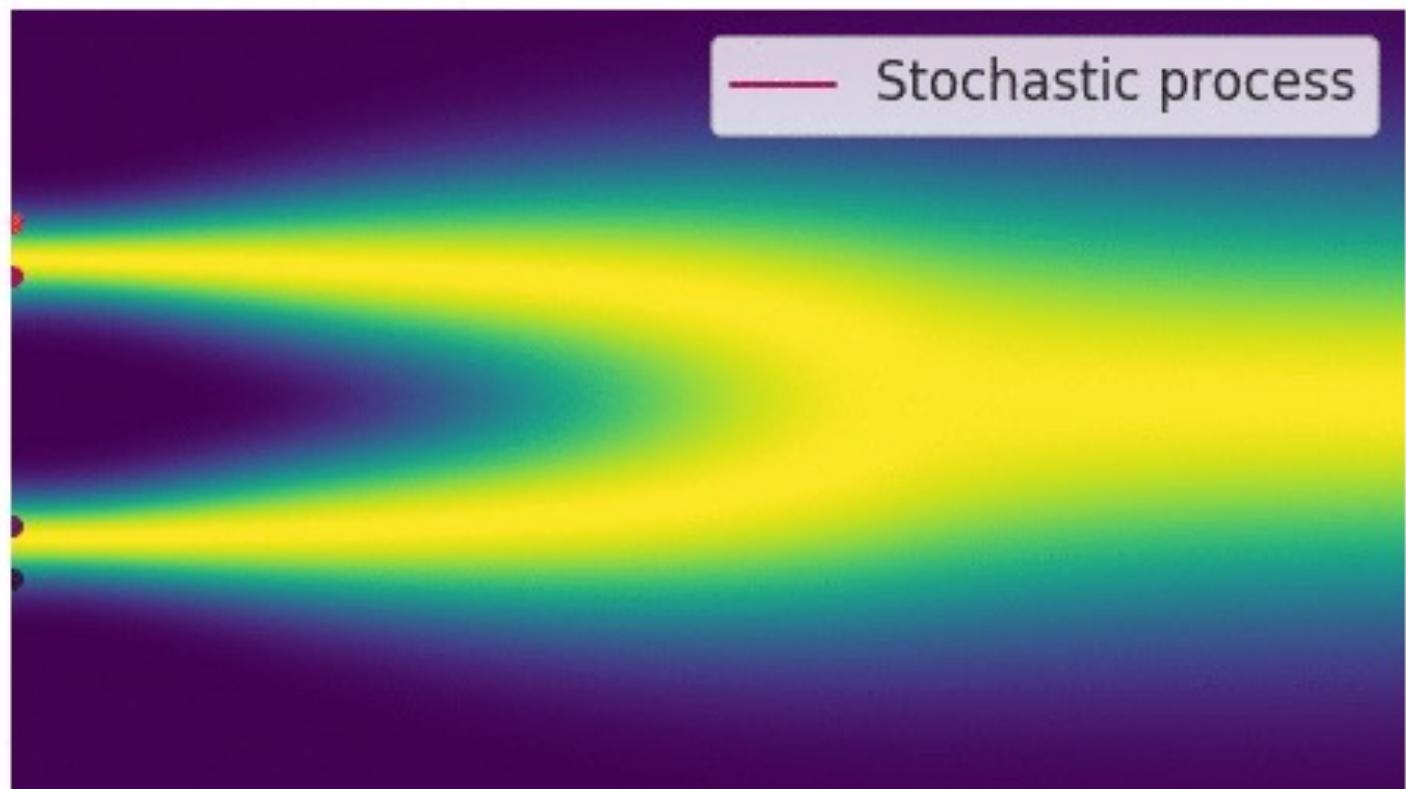
$$\mathbb{E}_{t \in \mathcal{U}(0,T)} \mathbb{E}_{p_t(\mathbf{x})} [\lambda(t) \|\nabla_{\mathbf{x}} \log p_t(\mathbf{x}) - \mathbf{s}_\theta(\mathbf{x}, t)\|_2^2],$$

$$\begin{aligned}\Delta \mathbf{x} &\leftarrow [\mathbf{f}(\mathbf{x}, t) - g^2(t) \mathbf{s}_\theta(\mathbf{x}, t)] \Delta t + g(t) \sqrt{|\Delta t|} \mathbf{z}_t \\ \mathbf{x} &\leftarrow \mathbf{x} + \Delta \mathbf{x} \\ t &\leftarrow t + \Delta t,\end{aligned}$$

# Denoising Score Matching



# Denoising Score Matching



# Conditional Flow Matching

**Diffusion  
Process**

**Closed-form  
conditional probability**

$$d\mathbf{x}_t = f_t(\mathbf{x}_t)dt + g_t(\mathbf{x}_t)d\mathbf{w} \longrightarrow p(\mathbf{x}_t | \mathbf{x}_0)$$

**Training**

$$\left\| s_\theta(\mathbf{x}_t) - \nabla \log p_t(\mathbf{x}_t | \mathbf{x}_0) \right\|^2$$

**Sampling**

$$\dot{\mathbf{x}}_t = f_t(\mathbf{x}_t) - \frac{1}{2}g_t^2 s_\theta(\mathbf{x}_t)$$

# Key Equations for Score-based Modelling

"Forward" SDE

$$d\vec{x}_t = f(\vec{x}_t, t) dt + g(t) dW \quad \text{with} \quad \vec{x}_0 \sim p_1$$

"Reverse" SDE

$$d\overleftarrow{x}_t = (f(\overleftarrow{x}_t, t) - g(t)^2 \nabla \log p_t(\overleftarrow{x}_t)) dt + g(t) dW \quad \text{with} \quad \overleftarrow{x}_0 \sim p_0$$

Probability flow ODE

$$dx_t = \left[ f(x_t, t) - \frac{1}{2} g(t)^2 \nabla \log p_t(x_t) \right] dt \quad \text{with} \quad x_0 \sim p_0$$

# Score-based generative modelling

*Training:* learn score  $\mathbf{s}_t(x) := \nabla \log p_t(x)$  by

1. Sample from target  $x_1 \sim p_1$
2. Run "forward" SDE to "noise"  $x_1$  until it becomes "simple"  $x_0 \sim p_0$
3. Minimize some loss  $\propto \|\hat{s}_t(x) - \mathbf{s}_t(x)\|^2$

*Inference:* sample using "reverse" SDE or prob-flow ODE

---

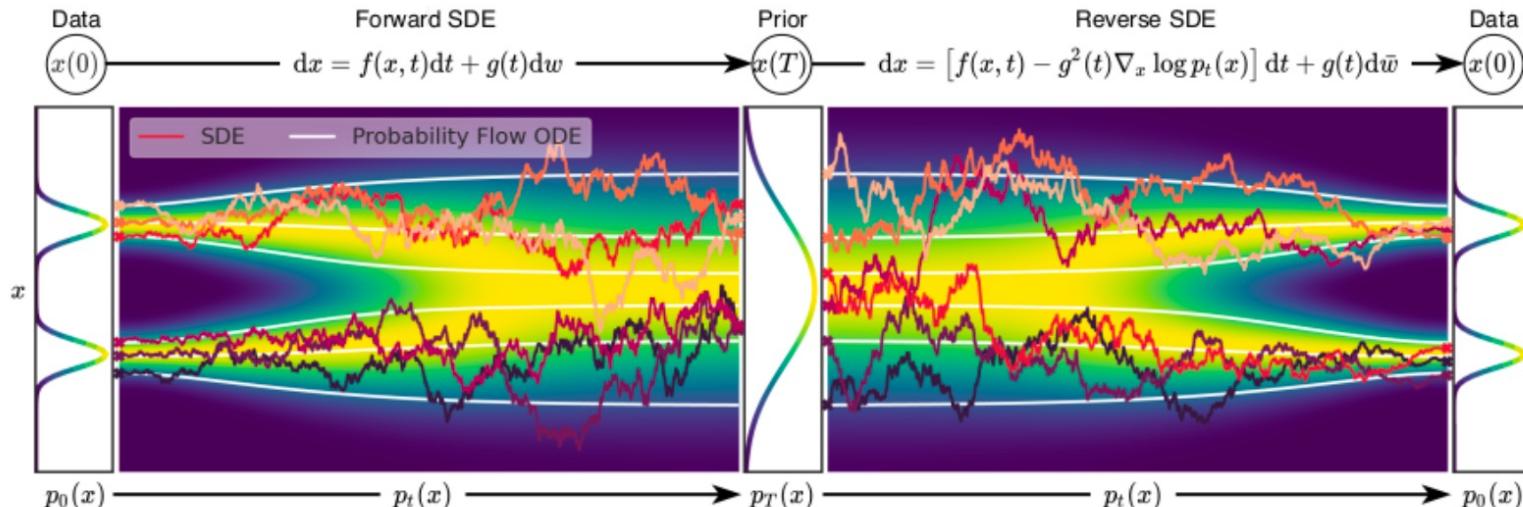


Figure 2 from Song et al. (2020)

# Conditional Flow Matching

**Diffusion  
Process**

**Closed-form  
conditional probability**

$$d\mathbf{x}_t = f_t(\mathbf{x}_t)dt + g_t(\mathbf{x}_t)d\mathbf{w} \longrightarrow p(\mathbf{x}_t | \mathbf{x}_0)$$

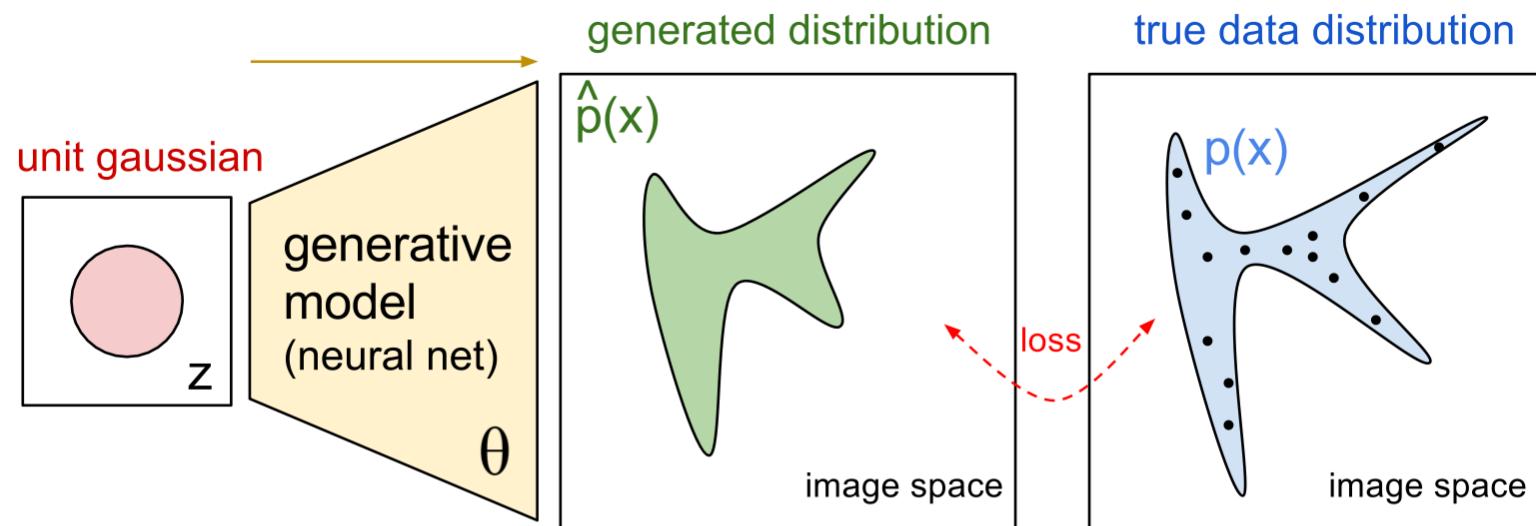
**Training**

$$\left\| s_\theta(\mathbf{x}_t) - \nabla \log p_t(\mathbf{x}_t | \mathbf{x}_0) \right\|^2$$

**Sampling**

$$\dot{\mathbf{x}}_t = f_t(\mathbf{x}_t) - \frac{1}{2}g_t^2 s_\theta(\mathbf{x}_t)$$

# Normalising Flows



Let  $p_0$  be a simple distribution on  $\mathbb{R}^d$  and  $\phi : \mathbb{R}^d \rightarrow \mathbb{R}^d$  a transformation (diffeomorphism).

Let  $p_1$  be the distribution of moving the samples of  $p_0$  along  $x_1 = \phi(x_0)$ .

$$p_1(x_1) = p_0(x_0) \left| \frac{\partial \phi}{\partial x_0}(x_0) \right|^{-1}, \quad \text{where } x_0 = \phi^{-1}(x_1)$$

# Normalising Flows

Parameterise the transformation by a deep neural network  $\hat{\phi}$ .

Maximum log likelihood objective:

$$\mathbb{E}_{x \sim \mathcal{D}} [\log p_1(x)] = \mathbb{E} \left[ \log p_0(x_0) - \log \left| \frac{\partial \hat{\phi}}{\partial x_0}(x_0) \right| \right], \quad \text{where } x_0 = \hat{\phi}^{-1}(x)$$

**Challenges** Requires computation of inverse  $\hat{\phi}^{-1}$  and the determinant of Jacobian  $\left| \frac{\partial \phi}{\partial z} \right|$

# Continuous Normalising Flows (CNFs)

Use multiple 'residual' transformations

$$x = (u_N \circ u_{N-1} \dots \circ u_1)(x_0).$$

For  $N \rightarrow \infty$ , the flow  $\phi_t$  describes the position of a starting point  $x_0$  along the vector field  $u_t$ , defined via an ODE

$$\frac{dx_t}{dt} = u_t(x_t),$$

where  $x_t = \phi_t(x_0)$  and  $u_t : \mathbb{R}^d \rightarrow \mathbb{R}^d$  is the tangent field of the flow.

The transformation is then the solution to

$$x \triangleq \phi_1(x_0) = x_0 + \int_0^1 u_t(x_t) dt.$$

# Continuous change-in-variables

Fokker-Planck equation without the diffusion term:

$$\log p_t(x) = \log p_0(x_0) - \int_0^t (\nabla \cdot u_t)(x_t) dt$$

Maximum likelihood training of Continuous Normalising Flows require:

- expensive numerical ODE simulations
- estimators for the divergence.

# Flow Matching

Integral-free approach to training CNF models

Supervised regression objective

Let  $u_t$  be a vector field that generates  $p_t$ , then we parameterise a neural net  $\hat{u} : \mathbb{R}_+ \times \mathbb{R}^d \rightarrow \mathbb{R}^d$  and want to learn it using

$$\mathcal{L} = \mathbb{E}_{t,p_t(x)} [\|\hat{u}(t, x) - u_t(x)\|^2]$$

**Challenges:**

How do we ensure  $p_1 \approx p$

What should  $p_0$  be?

What is  $u_t$ ?

Many names – same idea

## FLOW MATCHING FOR GENERATIVE MODELING

**Yaron Lipman<sup>1,2</sup> Ricky T. Q. Chen<sup>1</sup> Heli Ben-Hamu<sup>2</sup> Maximilian Nickel<sup>1</sup> Matt Le<sup>1</sup>**

<sup>1</sup>Meta AI (FAIR) <sup>2</sup>Weizmann Institute of Science

# Many names – same idea

## FLOW MATCHING FOR GENERATIVE MODELING

**Yaron Lipman<sup>1,2</sup> Ricky T. Q. Chen<sup>1</sup> Heli Ben-Hamu<sup>2</sup> Maximilian Nickel<sup>1</sup> Matt Le<sup>1</sup>**

<sup>1</sup>Meta AI (FAIR) <sup>2</sup>Weizmann Institute of Science

## BUILDING NORMALIZING FLOWS WITH STOCHASTIC INTERPOLANTS

**Michael S. Albergo**

Center for Cosmology and Particle Physics  
New York University  
New York, NY 10003, USA  
[albergo@nyu.edu](mailto:albergo@nyu.edu)

**Eric Vanden-Eijnden**

Courant Institute of Mathematical Sciences  
New York University  
New York, NY 10012, USA  
[eve2@cims.nyu.edu](mailto:eve2@cims.nyu.edu)

# Many names – same idea

FLOW MATCHING

Learning to Generate and Transfer Data with Rectified Flow

**Yaron Lipman<sup>1,2</sup>** **Ricky T. Q.**  
<sup>1</sup>Meta AI (FAIR) <sup>2</sup>Weizmann I

BUILDING NORMA  
INTERPOLANTS

**Michael S. Albergo**  
Center for Cosmology and Particle Physics  
New York University  
New York, NY 10003, USA  
albergo@nyu.edu

Xingchao Liu\*  
University of Texas at Austin  
xcliu@utexas.edu

Chengyue Gong\*  
University of Texas at Austin  
cygong@cs.utexas.edu

Qiang Liu  
University of Texas at Austin  
lqiang@cs.utexas.edu  
~~ERIC VANDEN-EIJNDWIJN~~  
Courant Institute of Mathematical Sciences  
New York University  
New York, NY 10012, USA  
eve2@cims.nyu.edu

# Many names – same idea

FLOW MATCHING

Learning to Generate and Transfer Data with Rectified Flow

**Yaron Lipman<sup>1,2</sup>** **Ricky T. Q.**  
<sup>1</sup>Meta AI (FAIR) <sup>2</sup>Weizmann I

BUILDING NORMA  
INTERPOLANTS

**Michael S. Albergo**  
Center for Cosmology and Particle Physics  
New York University  
New York, NY 10003, USA  
albergo@nyu.edu

Xingchao Liu\*  
University of Texas at Austin  
xcliu@utexas.edu

Chengyue Gong\*  
University of Texas at Austin  
cygong@cs.utexas.edu

Qiang Liu  
University of Texas at Austin  
lqliang@cs.utexas.edu

Action Matching:  
Learning Stochastic Dynamics from Samples

---

Kirill Neklyudov <sup>1</sup> Rob Brekelmans <sup>1</sup> Daniel Severo <sup>1,2</sup> Alireza Makhzani <sup>1,2</sup>

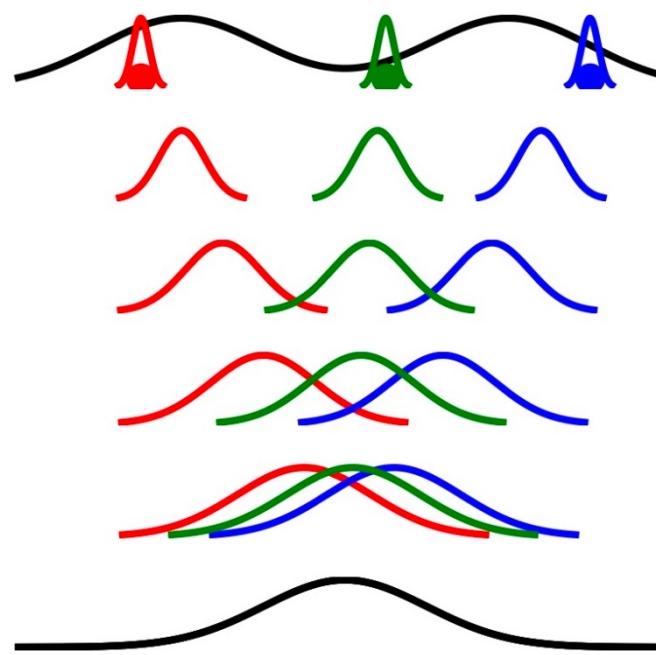
# Conditional Flow Matching

**Target probability path:** as mixture of simpler probability paths  
 $p_t = \int p_t(\cdot|x_1)p(x_1)dx_1$

**Conditional probability path**  $p_t(\cdot|x_1)$  s.t.  $p_1(\cdot|x_1) = \delta_{x_1}$  and  $p_0(\cdot|x_1) = p_0$

**Recover data distribution**

$$p_1(x) = \int p_1(x|x_1)q(x_1)dx_1 = \int \delta_{x_1}(x)q(x_1)dx_1 = q_1(x)$$

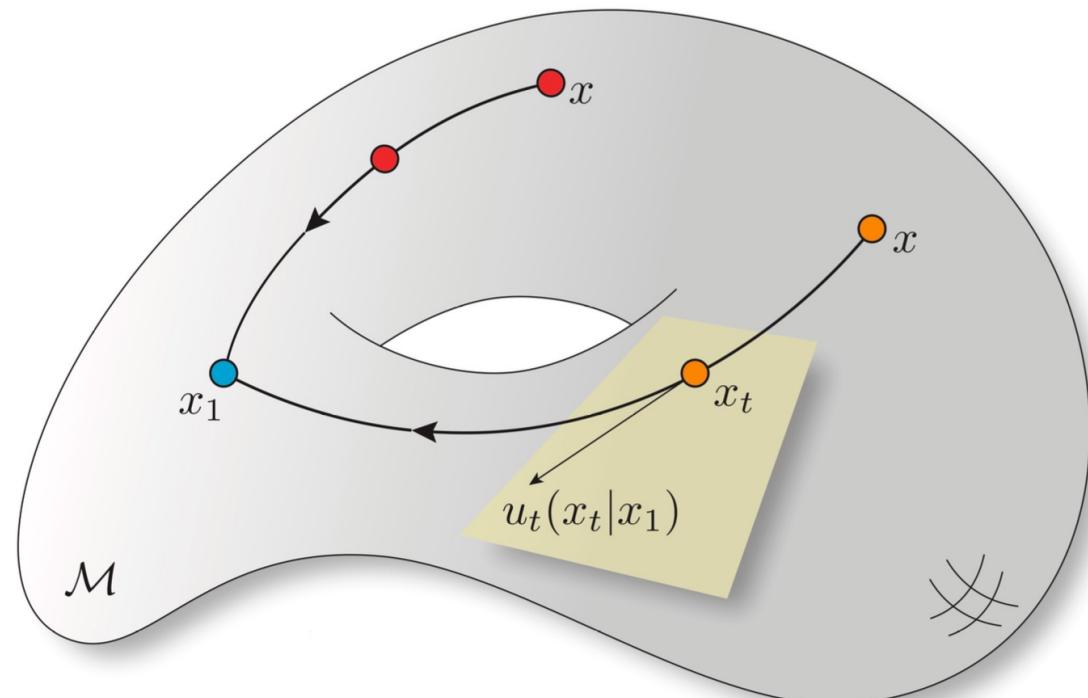


# Conditional Flows

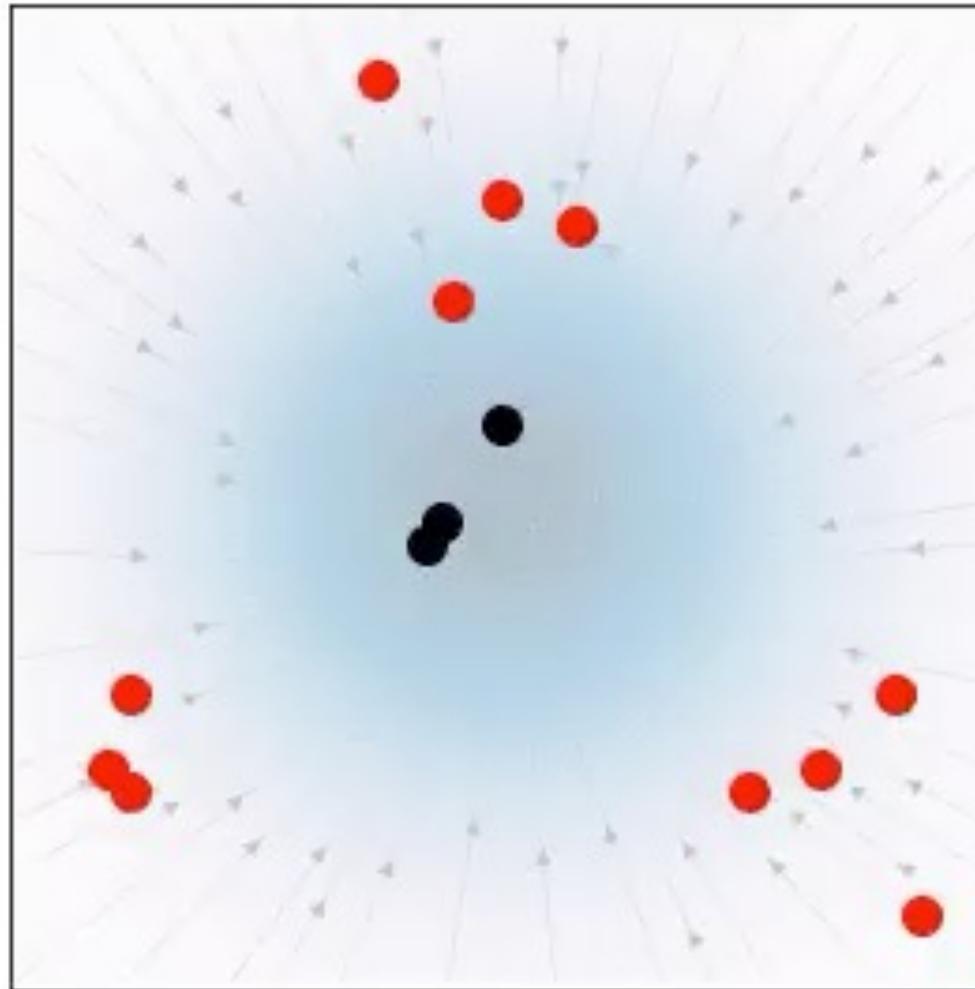
**Conditional vector field**  $u_t(\cdot|x_1)$  inducing conditional probability path  
 $p_t(\cdot|x_1)$

**Marginal vector field**  $u_t$  via *marginalising* over the *conditional* vector field  
 $u_t(\cdot|x_1)$

$$u_t(x) = \int u_t(x|x_1)p_t(x_1|x)dx_1 = \int u_t(x|x_1) \frac{p_t(x|x_1)q(x_1)}{p_t(x)} dx_1 = \mathbb{E}_{x_0, x_t \sim q_1 p_{t|1}}[u_t(x_t|x_1)]$$



# Conditional Flows



# (Conditional) Flow Matching: Training

**(Exact) flow matching:**  $\mathcal{L}_{\text{FM}}(\theta) = \mathbb{E}_{tx_t \sim \mathcal{U}[0,1]p_t} [\|u_\theta(t, x_t) - u_t(x_t)\|^2]$  with  
 $u_t(x) = \mathbb{E}_{x_0, x_t \sim q_1 p_{t|1}} [u_t(x_t|x_1)]$

Akin to score matching, one can actually move the expectation outside the  $\ell^2$  norm

**Conditional flow matching:**

$$\mathcal{L}_{\text{CFM}}(\theta) = \mathbb{E}_{t, x_0, x_t \sim \mathcal{U}[0,1]q_1 p_{t|1}} [\|u_\theta(t, x_t) - u_t(x_t|x_1)\|^2]$$

$$\nabla_\theta \mathcal{L}_{\text{FM}}(\theta) = \nabla_\theta \mathcal{L}_{\text{CFM}}(\theta)$$

Sampling  $x_t$  and evaluating  $u(x_t|x_1)$  is available in closed form.

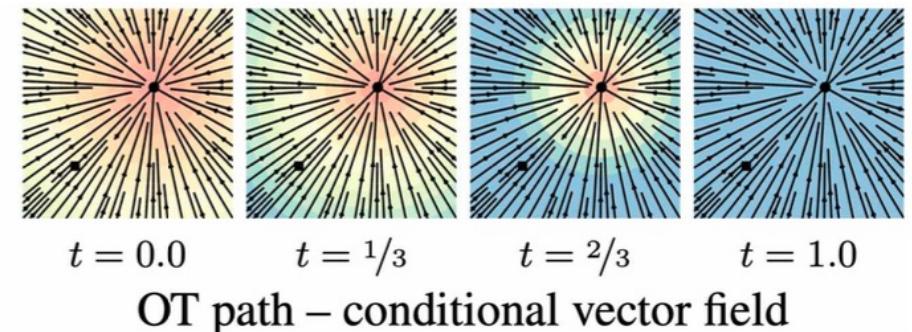
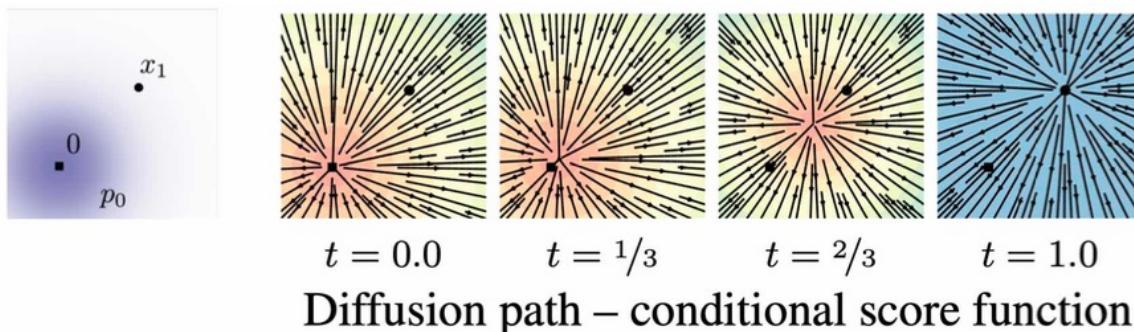
# Gaussian Probability Paths

**Conditional Probability path:**  $p_t(x|x_1) = \mathcal{N}(\mu_t(x_1), \sigma_t(x_1)^2 \mathbf{I})$

**Conditional vector field:**  $u_t(x|x_1) = \frac{\sigma'_t(x_1)}{\sigma_t(x_1)}(x - \mu_t(x_1)) + \mu'_t(x_1)$

**Example: Linear interpolation**

- $\mu_t \triangleq tx_1$  and  $\sigma_t \triangleq 1 - t \Rightarrow p_t(x|x_1) = \mathcal{N}(x|tx_1, (1-t)^2)$
- $u_t(x|x_1) = \frac{1}{1-t}(x_1 - x)$



# Flow Matching vs. Diffusion

---

**Algorithm 1:** Flow Matching training.

---

**Input :** dataset  $q$ , noise  $p$

Initialize  $v^\theta$

**while** not converged **do**

$t \sim \mathcal{U}([0, 1])$	▷ sample time
$x_1 \sim q(x_1)$	▷ sample data
$x_0 \sim p(x_0)$	▷ sample noise
$x_t = \Psi_t(x_0   x_1)$	▷ conditional flow

Gradient step with  $\nabla_\theta \|v_t^\theta(x_t) - \dot{x}_t\|^2$

**Output:**  $v^\theta$

---

$p_t(x_t | x_1)$  general

$p(x_0)$  is general

---

**Algorithm 2:** Diffusion training.

---

**Input :** dataset  $q$ , noise  $p$

Initialize  $s^\theta$

**while** not converged **do**

$t \sim \mathcal{U}([0, 1])$	▷ sample time
$x_1 \sim q(x_1)$	▷ sample data
$x_t = p_t(x_t   x_1)$	▷ sample conditional prob

Gradient step with  
 $\nabla_\theta \|s_t^\theta(x_t) - \nabla_{x_t} \log p_t(x_t | x_1)\|^2$

**Output:**  $v^\theta$

---

$p_t(x_t | x_1)$  closed-form from of SDE  $dx_t = f_t dt + g_t dw$

- **Variance Exploding:**  $p_t(x | x_1) = \mathcal{N}(x | x_1, \sigma_{1-t}^2 I)$
- **Variance Preserving:**  $p_t(x | x_1) = \mathcal{N}(x | \alpha_{1-t} x_1, (1 - \alpha_{1-t}^2) I)$   
 $\alpha_t = e^{-\frac{1}{2}T(t)}$

$p(x_0)$  is Gaussian

$p_0(\cdot | x_1) \approx p$

# Conditional Flow Matching

**"So, what's the difference between FMs and SBDMs?"**

	<b>Learn</b>	<b>ODE inference</b>	<b>SDE inference</b>	<b>Exact endpoints</b>
FM	$u(t, x)$	✓	✗ <sup>1</sup>	✓
SBDM	$s(t, x)$	✓	✓	✗ <sup>2</sup>

Which is "better"? 🤔

Note: unclear whether "exact endpoint" and "normalisable guarantee" matters in practice

# Score-based modelling

**Diffusion  
Process**

$$d\mathbf{x}_t = f_t(\mathbf{x}_t)dt + g_t(\mathbf{x}_t)d\mathbf{w} \longrightarrow p(\mathbf{x}_t | \mathbf{x}_0)$$

**Closed-form  
conditional probability**

**Training**

$$\left\| s_\theta(\mathbf{x}_t) - \nabla \log p_t(\mathbf{x}_t | \mathbf{x}_0) \right\|^2$$

**Sampling**

$$\dot{\mathbf{x}}_t = f_t(\mathbf{x}_t) - \frac{1}{2}g_t^2 s_\theta(\mathbf{x}_t)$$

# Conditional Flow Matching

**General conditional  
probability path**

$$p(\mathbf{x}_t | \mathbf{x}_0) = \mathcal{N}(\mu_t(\mathbf{x}_0), \sigma_t^2(\mathbf{x}_0)I)$$

**Training**

$$\left\| v_{\theta}(\mathbf{x}_t) - u_t(\mathbf{x}_t | \mathbf{x}_0) \right\|^2$$

**Sampling**

$$\dot{\mathbf{x}}_t = v_{\theta}(\mathbf{x}_t)$$